



Vault Enterprise Technical Overview & Architectural Deep-Dive

February 2023

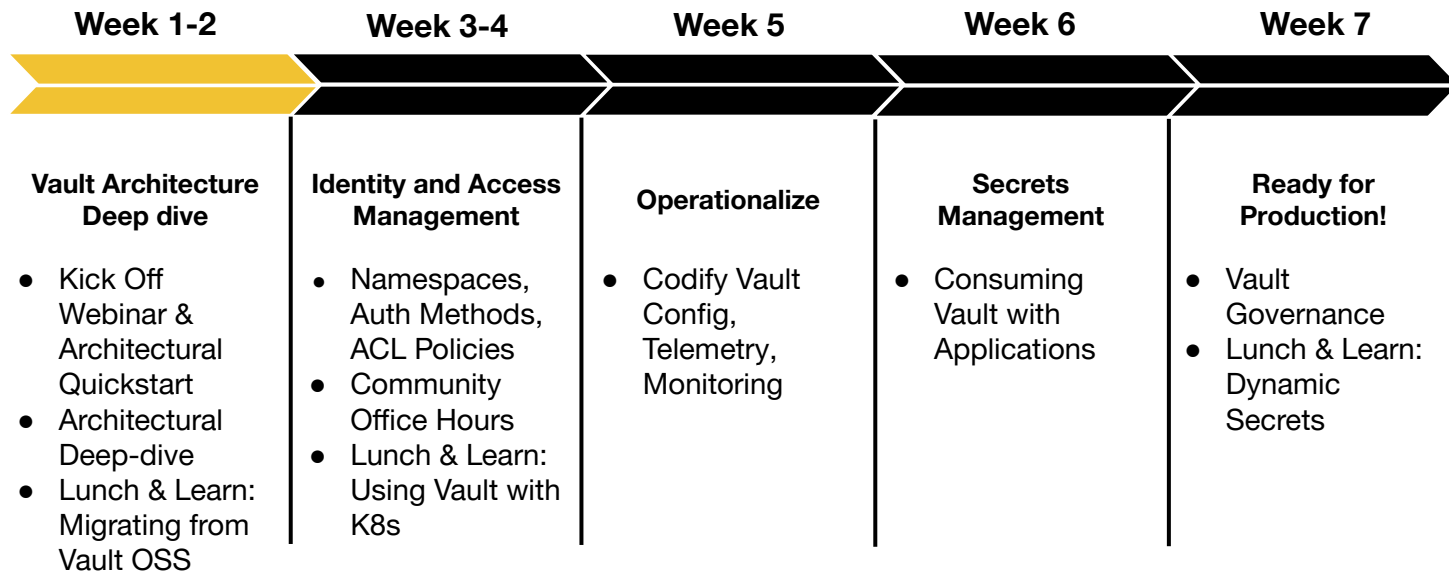
Copyright © 2021 HashiCorp



Agenda

1. Overview
2. Architecture
3. Deployment Patterns
4. Operations
5. Next Steps

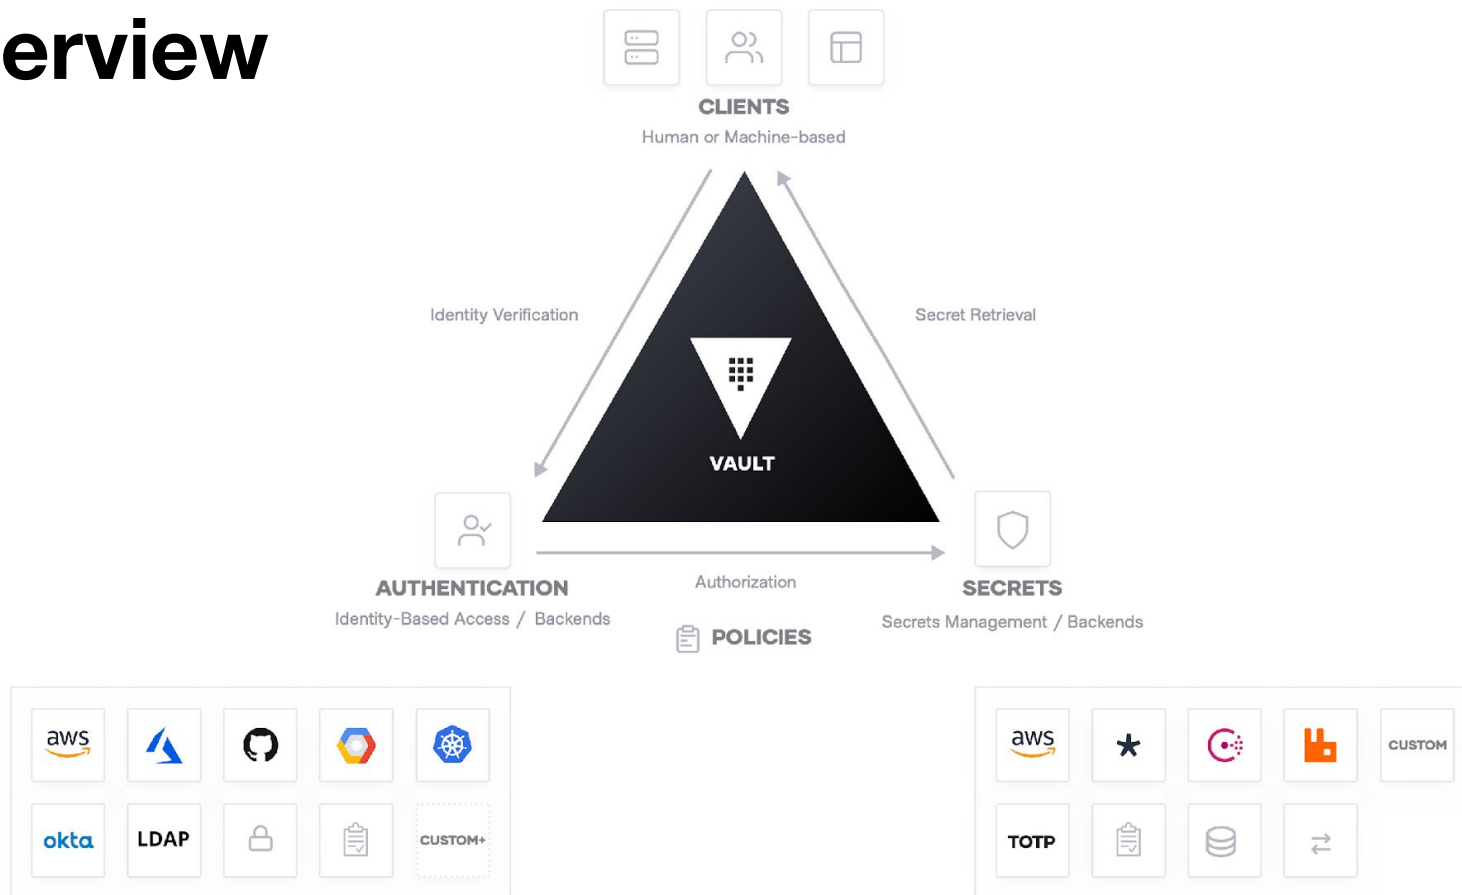
Vault Enterprise Path to Production



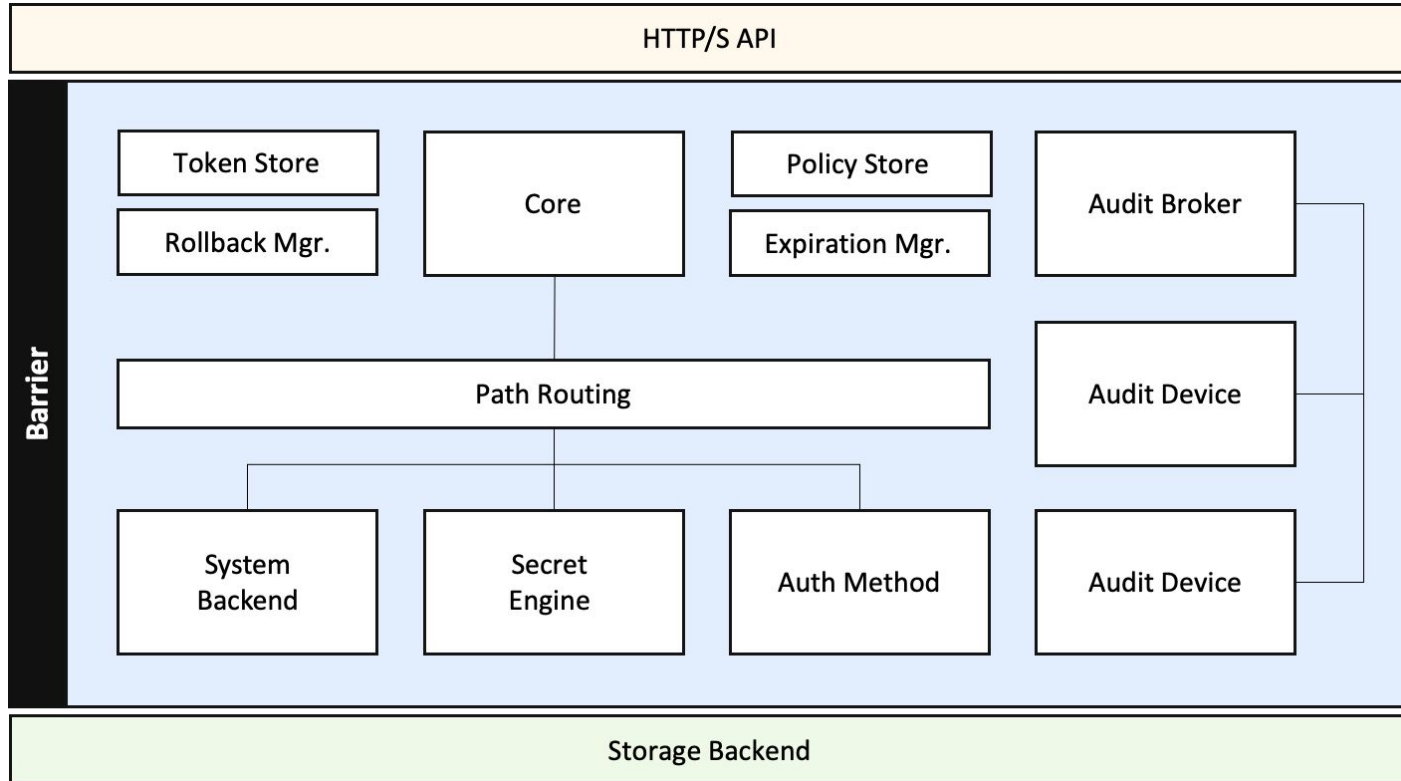
— 01

Vault Overview

Overview



Architecture & Cryptographic Barrier



Vault Security Model



- It's all about access to the Encryption Key
- Configuring `cap_ipc_lock=+ep`, `LimitNOFILE`, and `LimitMEMLOCK` prevent Memory Swapping to Disk, so secrets are not written in plain text to disk
- The Vault Encryption Key is stored in memory in **PLAIN TEXT**
 - This is done for performance
 - Root access to an unlocked vault server could compromise this
 - Isolation technologies which allow reading of memory could compromise this (VM issues, but principally Kubernetes)
- Master Root Key protects the Encryption key, so it also must be secure
- Auto-Unseal is a recommended pattern as it shifts the risk profile

Cryptography Security Model



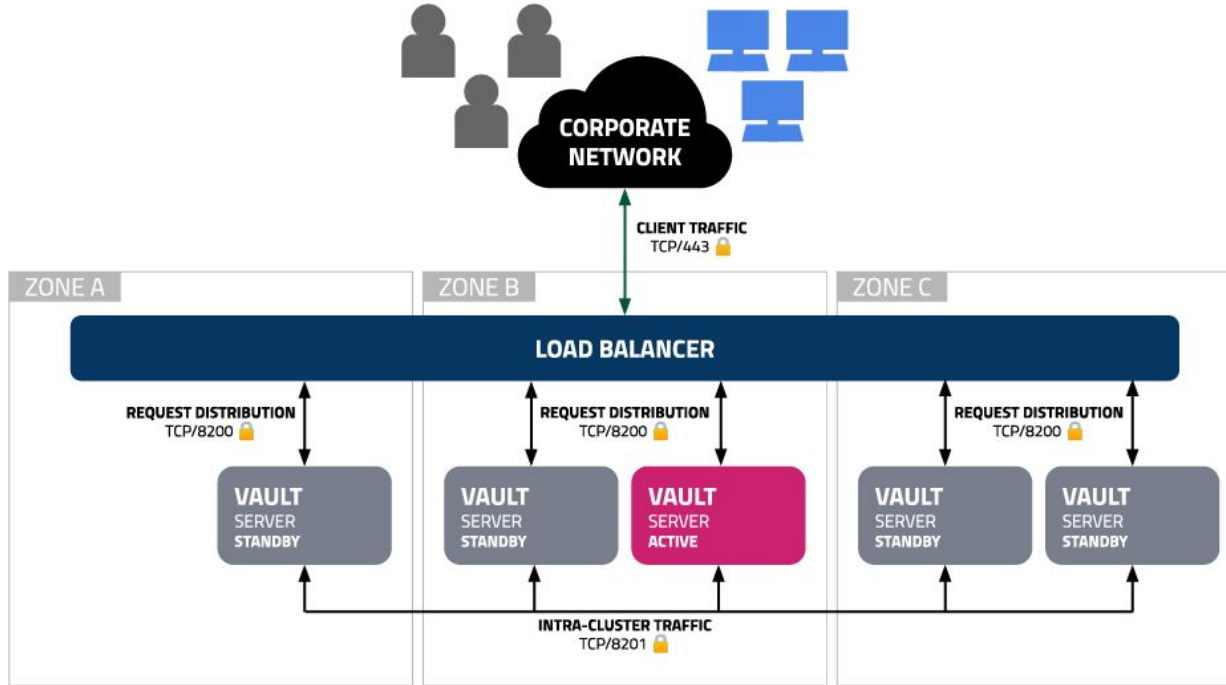
- **Vault uses publicly available cryptographic technologies**
- P vs NP - Good cryptographic algorithms are exponential in difficulty to solve but polynomial in difficulty to validate answers for
- Numerous algorithms (SHA1) were exposed to have defects that allowed them, or a subset of them to be reduced to polynomial difficulty problems
- Short encryption keys and faster computers has made brute-forcing older encryption standards possible
- Software based random number generations suffer from a lack of randomness

— 02

Vault Architecture



Integrated Storage Reference Architecture



5 Vault Servers across 3 Availability Zones

Vault Integrated Storage Architecture



- **Integrated Storage Autopilot**

- Monitors node health status
- Server stabilization - prevent quorum disruption from an unstable node
- Dead server cleanup
- Enabled by default in Vault 1.7.0 and higher

- **Vault 1.11.0+ new features**

- Automated upgrades promotes new versioned nodes to voter nodes removed old versioned nodes
- Redundancy zones allows for deployment of non-voter nodes in an AZ with automatic promotion if a node is lost



Sizing

Per instance sizing
recommendations

	Small (Dev/Test/Staging/QA)	Large (Production)
CPU	2 - 4 Core	8 - 16 Core
Memory	8 - 16 GB RAM	32 - 64 GB RAM
Disk Capacity	100+ GB	200+ GB
Disk IO	3000+ IOPS	7500+ IOPS
Disk Throughput	75+ MB/s	250+ MB/s
AWS	m5.large, m5.xlarge	m5.2xlarge, m5.4xlarge
Azure	standard_d2s_v3, standard_d4s_v3	standard_d8s_v3, standard_d16s_v3
GCP	n2-standard-2, n2-standard-4	n2-standard-8, n2-standard-16

Performance Considerations



Profile Workloads

- As Vault adoption scales throughout an organization there will be varying workloads utilizing Vault
- Different workloads have varying impacts to resources (RAM, CPU, I/O)
- Leverage telemetry monitoring to ensure an understanding of implications to Vault Cluster resources usage
- As new applications/services/teams/users are onboarded to Vault, profile the usage patterns to ensure optimal authentication and consumption patterns are used

Performance Considerations



External Systems

- Authentication Methods & Secrets Engines have external systems dependencies that can impact Vault's ability to process requests
- Ensure telemetry is enabled on those systems and services and proactively monitor for performance issues

Networking Considerations



Integrated Storage is network latency dependent

- <8ms RT network connection required to ensure Raft Storage remains consistent across all Vault Nodes.
- Restrict communication to only required ports and CIDRs
- Standard HTTPS TLS encryption should be used to protect network traffic



Networking Requirements



Source	Destination	Port	Protocol	Direction	Purpose
Client Machines	Load Balancer	443	tcp	incoming	Request distribution
Load Balancer	Vault Servers	8200	tcp	incoming	Vault API
Vault Servers	Vault Servers	8200	tcp	bidirectional	Cluster Bootstrapping
Vault Servers	Vault Servers	8201	tcp	bidirectional	Raft, replication, request forwarding
Vault Servers	External Systems	various	various	various	External APIs

Load Balancing



Vault does not include built in load balancing capabilities

- To ensure Vault availability and reliability either an external load balancer or Consul should be used to distribute client requests
- **TLS should terminate at Vault** and not the load balancer to ensure end-to-end encryption
- Use Vault's health endpoint to determine active node and node health `https://<vaultnode>:8200/v1/sys/health`



Scaling Considerations

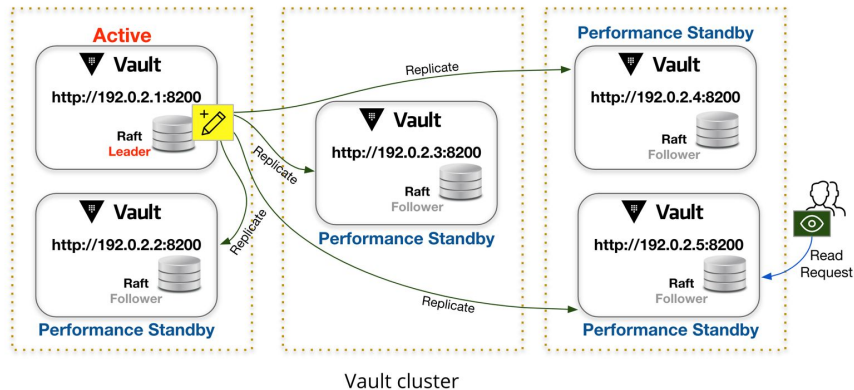


Managed scaling services should be leveraged when deploying in a cloud environment to ensure the Vault cluster remains populated with health nodes

- Additional nodes will not increase performance
- Do not replace all instances at once in a scaling group otherwise data-loss will occur

Cloud	Managed Auto Scaling Service
AWS	Auto Scaling Group
Azure	Virtual Machine Scale Sets
GCP	Managed Instance Groups

Scaling Performance Standby Nodes



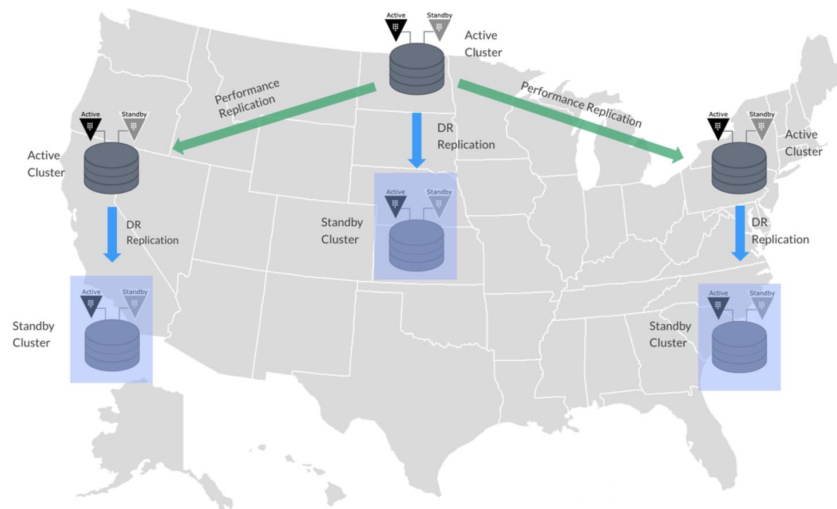
Horizontal scalability for read requests

- Performance Standby Nodes can be used to respond to read-only requests
- Performance Standby Nodes are enabled by default and process read-only requests locally
- Write requests are forwarded to the Active Node
- Integrated Storage uses eventual consistency and data may not be available across all nodes immediately
- Vault 1.7+ includes multiple methods to control how requests are handled

Vault Replication



- Vault can be extended to multiple regions using replication
- The primary cluster uses asynchronous replication to ship data to the secondaries
- Multiple replication modes can be combined to provide resilience and performance



Replication Types




Disaster Recovery

- Provides a warm standby cluster that contains all data from the primary Vault cluster
- **It is strongly recommended to deploy at least one DR cluster**

Performance Replication

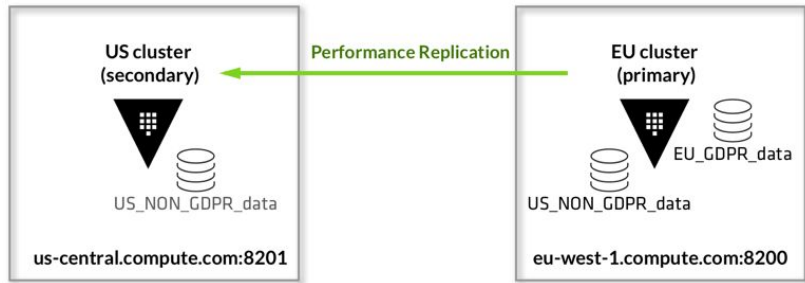
- Provides an active Vault cluster with shared state of the primary
- Replicates: secrets, authentication methods, policies, & other shared data
- Token and leases are not replicated to performance secondaries

Vault Replication



Disaster Recovery (DR) Replication

- Achieve RPO/RTO requirements
- Vault is typically considered a Tier 0 application



Performance Replication

- Additional cluster closer to source of requests
- Latency reduction, compliance and data sovereignty
- Segment certain types of workloads

03

Deployment Patterns

Recommended Patterns



Immutable Builds

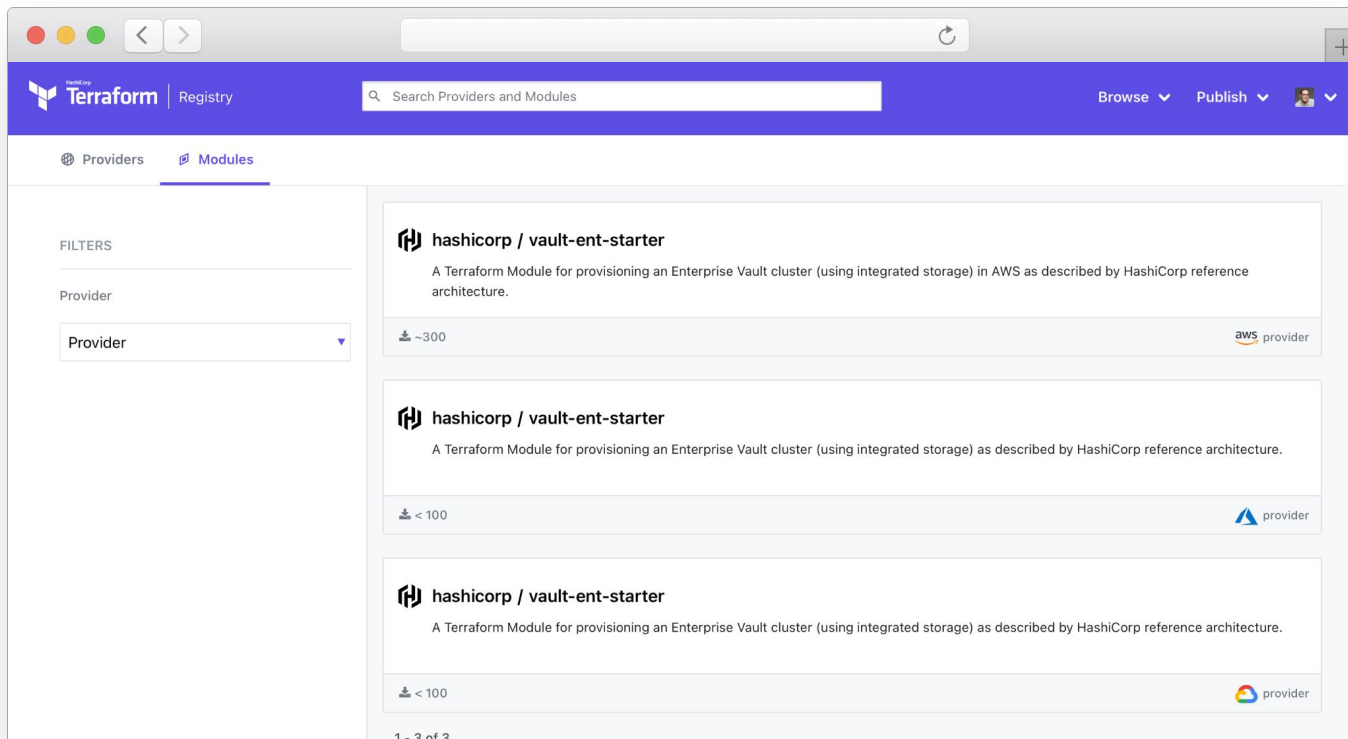
- Tools like Packer can be used to build immutable machine for blue/green deployment using existing CI/CD orchestration
- This can streamline lifecycle processes for managing Vault
- When using this pattern with Integrated Storage, ensure measures are taken to ensure quorum is maintained as new image versions are introduced to the cluster

Configuration Management

- Configuration Management tools and patterns can be used for installation, upgrade, and configuration of Consul
- Autopilot can be leveraged for in-place upgrades (Vault 1.11.0+)

Terraform Modules

Quickly deploy Vault cluster based on reference architecture - [link](#)



Vault Helm Chart

Deploy Vault Reference Architecture inside Kubernetes - [link](#)



The screenshot shows the GitHub repository page for `hashicorp/vault-helm`. The repository is public and has 50 watches, 622 stars, and 515 forks. The main branch is `main`, with 4 branches and 25 tags. The repository contains a `fix chart publish job (#620)` commit, which is the latest release. The repository also includes a `Readme`, `MPL-2.0 License`, and 24 releases. The latest release is `v0.16.1`, which was published 9 days ago. The repository also includes a `Chart.yaml` file, which is the main configuration file for the Helm chart.

File	Commit	Time
<code>fix chart publish job (#620)</code>	<code>9758666</code>	9 days ago
<code>fix chart publish job (#620)</code>	<code>9758666</code>	9 days ago
<code>Update jira sync github action (#411)</code>	<code>9758666</code>	11 months ago
<code>Adding support for the old leader-elect</code>	<code>9758666</code>	23 days ago
<code>vault-helm 0.16.1 release</code>	<code>9758666</code>	9 days ago
<code>feature: Support configuring various properties as YAML directly</code>	<code>9758666</code>	3 months ago
<code>Ignore bin dirs</code>	<code>9758666</code>	3 years ago
<code>vault-helm 0.16.1 release</code>	<code>9758666</code>	9 days ago
<code>vault-helm default branch is now main</code>	<code>9758666</code>	11 days ago
<code>vault-helm 0.16.1 release</code>	<code>9758666</code>	9 days ago



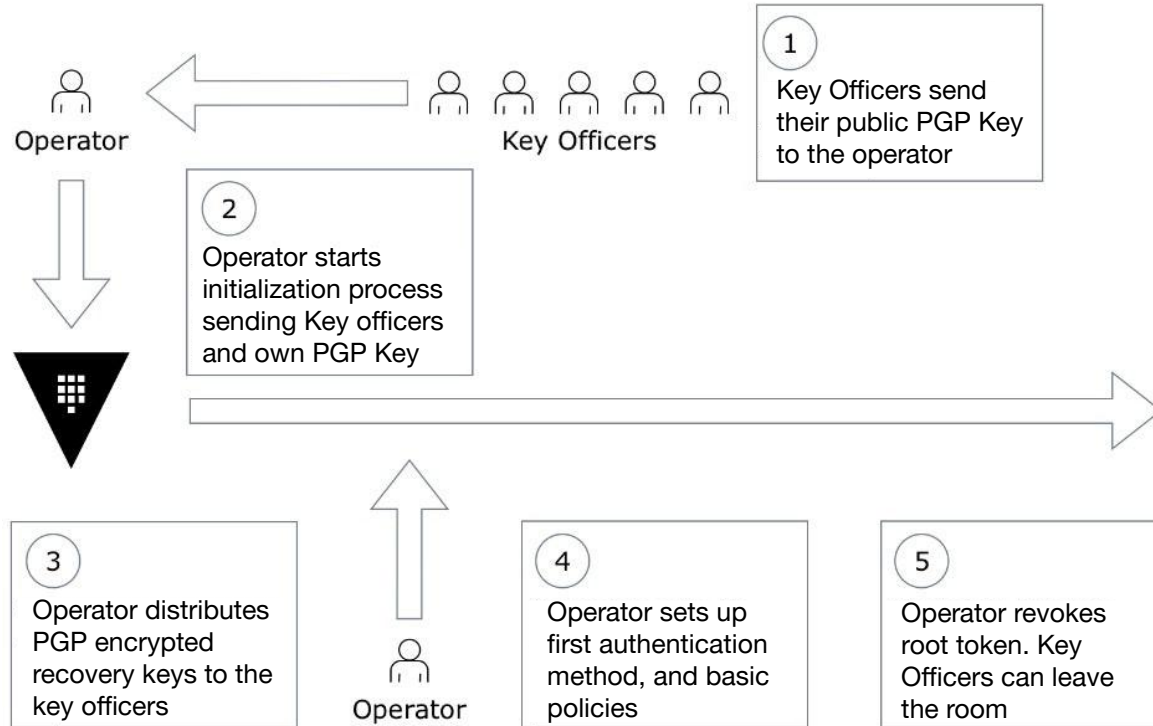
Upgrades

- Major upgrades should occur **at least 2X per year** to stay within **N-2 major releases** version support window
- Automation of the update process is recommended to ensure ease of operations and keep Consul patched with current updates
- Prior to a production upgrade:
 - Review [version specific upgrade guide](#)
 - Review [changelog](#)
 - Test version in QA environment
 - Take a snapshot prior to any upgrade

— 04

Operations

Vault Cluster Initialization




Auto-unseal



Unsealing is the process of constructing the master key necessary to decrypt the data encryption key because by default, Vault needs to be unsealed before any operation can be performed

Vault supports auto-unseal from:


- HSM
- AliCloud KMS
- AWS KMS
- Azure Key Vault
- Google Cloud KMS
- OCI KMS



HSM Integration



- Integrate Vault with FIPS 140-2 certified HSM (Hardware Security Module) and enable the Seal Wrap feature to protect your data.
- Vault encrypts secrets using 256-bit AES in GCM mode with a randomly generated nonce prior to writing them to its persistent storage. When you enable seal wrap, Vault wraps your secrets with an extra layer of encryption leveraging the HSM encryption and decryption.





Root Token Generation

TERMINAL

```
> vault operator init  
Unseal Key 1: Ly7wgNFzKVcw95nv6fLTQ/lsf49Wn4JaIEYGPm15pSzn  
Unseal Key 2: JWeteKjgpFXI2wY2I16j8JCCy92P04GxGCyXvLCoHp1L  
Unseal Key 3: zLkMb09Lcr3QRwIgw7KBPy5jRD9aUttt010HZ4dusvx  
Unseal Key 4: 0J5fD29c5ZisK11jL13K0XOmIWu66PfA6NBV3UEK7f/f  
Unseal Key 5: ahR01B203Kzxv0a0HgBLUDmByxhFdeVOVeA3l6PMIKMn
```



```
Initial Root Token: s.dZ1m130RBFkF0rQeWtLF3uiA
```

Vault initialized with 5 key shares and a key threshold of 3. Please securely distribute the key shares printed above. When the Vault is re-sealed, restarted, or stopped, you must supply at least 3 of these keys to unseal it before it can start servicing requests.

Vault does not store the generated master key. Without at least 3 key to reconstruct the master key, Vault will remain permanently sealed!

It is possible to generate new unseal keys, provided you have a quorum of existing unseal keys shares. See "vault operator rekey" for more information.



Root Token Handling Practices

The root token is returned to the operator during the initialization ceremony. This token can do **anything** in Vault and its usage should be closely monitored.

- Once operator has configured a secondary authentication method and granted administrators sudo access, almost all operations can be performed
- Best practice is **NOT** persisting the root token
- Generate a root token only when absolutely necessary



Production Readiness

Critical items to have in place
before production go-live

Backup

Automated Integrated Storage Snapshots, a Vault Enterprise feature takes periodic snapshots of Vault's data data

- Determine where snapshot files will be stored
- Configure based off your RPO/RTO requirements
- If snapshot is restored, the unseal keys that were valid at the time of the snapshot will be used to unseal



Automated Integrated Storage Snapshots

```
TERMINAL

> vault write \
    sys/storage/raft/snapshot-auto/config/testsnap \
    storage_type=local \
    file_prefix=testsnappy \
    interval=120m \
    retain=7 \
    local_max_space=1000000 \
    path_prefix=/opt/vault/
```



Production Readiness

Critical items to have in place
before production go-live

Monitoring

Vault should be monitored closely to ensure the service remains healthy and available in production

- Telemetry - Export telemetry data to solution that can analyze and identify trends overtime
- Log Analytics - Capture app logs and system logs and perform analysis on the log files for useful signals
- Active Health Checks - Query health endpoints to get the health of nodes and route traffic to active node



Production Readiness

Critical items to have in place before production go-live

Auditing

Vault sends audit information to a SIEM system or logging backend

- Determine audit devices that will be used
- Vault will not respond if the audit device is unavailable, use multiple audit devices to ensure Vault remains available
- Sensitive fields are HMAC, Selectively determine if any HMAC fields need to be exposed

— 05

Next Steps

Tutorials

<https://developer.hashicorp.com/vault/tutorials>



Step-by-step guides to accelerate deployment of Vault

The screenshot shows the HashiCorp Vault Developer Tutorials page. The top navigation bar includes links for Vault, Install, Tutorials, Documentation, API, Integrations, and Try Cloud. A search bar is located on the right. The left sidebar contains a 'Vault Home' link and a 'Tutorials' section with links for 'Get Started', 'CLI Quick Start', 'HCP Vault Quick Start', and 'UI Quick Start'. Below this are 'Use Cases' (ADP, Database Credentials, Data Encryption, Key Management, Secrets Management) and 'Certification Prep' (Associate). The main content area features a breadcrumb trail 'Developer / Vault / Tutorials' and a large yellow banner titled 'Centrally store, access and deploy secrets' with a link to 'Explore HCP Vault'. Below the banner is a 'Get Started' section with three cards: 'Getting Started' (13 tutorials), 'Getting Started with Vault UI' (8 tutorials), and 'Getting Started with HCP Vault' (9 tutorials). The 'New Tutorials' section lists 'Disaster Recovery Replication Failover and Failback' (54min) and 'Vault Installation to Minikube via Helm with TLS enabled' (13min).

Developer / Vault / Tutorials

Centrally store, access and deploy secrets

Explore HCP Vault →

Get Started

- 13 tutorials
Getting Started
Vault secures, stores, and tightly controls access to tokens, passwords, certificates, API keys,...
- 8 tutorials
Getting Started with Vault UI
Manage Vault environment as well as your secrets using Vault UI.
- 9 tutorials
Getting Started with HCP Vault
Quickly get hands-on with HashiCorp Cloud Platform (HCP) Vault using the HCP portal and...

New Tutorials

Here are the most recently published tutorials.

- 54min
Disaster Recovery Replication Failover and Failback
Learn how to failover in a Disaster Recovery Replication environment and then failback to original operating state.
- 13min
Vault Installation to Minikube via Helm with TLS enabled
Deploy Vault on Kubernetes locally with TLS using Minikube and the official Helm chart.

On this page

- Get Started
- New Tutorials
- HashiCorp Well-Architected ...
- Popular Topics
- All Tutorials



Resources

- [Vault Internal Architecture](#)
- [Vault Security Model](#)
- [Vault Reference Architecture](#)
- [Vault Redundancy Zones \(1.11.0+\)](#)
- [Terraform Starter Code](#)
- [Disaster Recovery Replication Setup](#)
- [Performance Replication Setup](#)
- [Vault Eventual Consistency and Controls](#)

Need Additional Help?



Customer Success

Contact our Customer Success Management team with any questions. We will help coordinate the right resources for you to get your questions answered.

customer.success@hashicorp.com

Technical Support

Something not working quite right? Engage with HashiCorp Technical Support by opening a ticket for your issue at support.hashicorp.com.

Discuss

Engage with the HashiCorp Cloud community including HashiCorp Architects and Engineers

discuss.hashicorp.com

Upcoming Webinars



Vault

Migrating from Vault OSS to Enterprise

This Lunch & Learn (separate link) covers the best methods for upgrading an existing Vault OSS Cluster to Vault Enterprise

Auth Methods, Namespaces & Policy

Learn best practices for deploying Vault Namespaces, Authentication Methods, and Vault policy

Using Vault with Kubernetes

This Lunch & Learn (separate link) covers the best practices for integrating Vault Enterprise with Kubernetes and

Action Items



- Share to customer.success@hashicorp.com
 - Authorized technical contacts for support
 - Stakeholders contact information (name and email addresses)
- Email krystal.allen@hashicorp.com with a brief summary of Vault Enterprise use case(s), goals, and project milestones
- Determine Vault pattern and begin deployment of first cluster(s)



Q & A





Thank You

customer.success@hashicorp.com

www.hashicorp.com